

Chapter 2

Memory Architecture for Multiprocessing





Outline

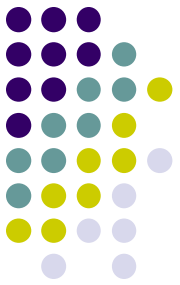
- Memory and Communication Issues
- Advanced Memory Hierarchy: Virtual Machines
- Multicore Memory Architecture: Commercial Examples



Memory and Communication Issues

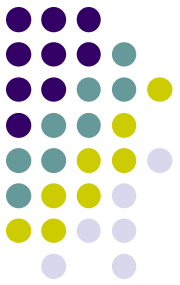
- ***Memory and Communication Issues***
 - Flynn's Taxonomy of Computer Categorization
 - Shared Memory Architecture: SMP/DSM
 - Shared Memory Communication
 - Cache Coherence in Shared Memory Architecture
 - Distributed Memory Architecture
 - Distributed Memory Communication: Message-Passing
 - Shared Memory vs. Distributed Memory
 - Memory Consistency
 - Parallel Computer Classification via Memory View
- Advance Memory Hierarchy: Virtual Machines
- Multicore Memory Architecture: Commercial Examples

Flynn's Taxonomy of Computer Categorization (1/2):



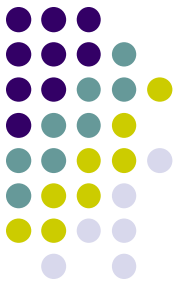
- Flynn's taxonomy is a classification of computer architectures, proposed by Michael J. Flynn in 1966
 - Single Instruction stream, Single Data stream (SISD)
 - e.g., traditional uniprocessor machines like a PC or old mainframes
 - Single Instruction stream, Multiple Data streams (SIMD)
 - e.g., an array processor, vector processing unit, or GPU
 - Multiple Instruction streams, Single Data stream (MISD)
 - Uncommon architecture; may be used for fault tolerance
 - A number of systems operate on the same data stream and must agree on the result. Examples include the Space Shuttle flight control computer
 - Multiple Instruction streams, Multiple Data streams (MIMD)
 - Multiple autonomous processors simultaneously executing different instructions on different data. Distributed systems are generally recognized to be MIMD architectures; either exploiting a single shared memory space or a distributed memory space

Flynn's Taxonomy of Computer Categorization (2/2):




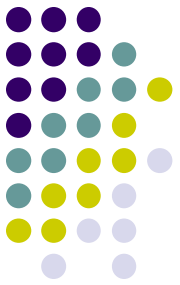
- (Following emphasize on parallel machines)
 - SISD machines
 - systems contain one CPU and hence can accommodate one instruction stream that is executed serially (conceptually)
 - **SIMD machines**
 - Exploits [data level parallelism](#)
 - Systems often have a large number of processing units
 - A single instruction manipulates many data items in parallel
 - A subclass of SIMD systems is vector processors
 - MISD machines
 - Theoretically, in this type of machines multiple instructions act on a single data stream
 - **MIMD machines**
 - Exploits [thread/program level parallelism](#)
 - Execute several instruction streams in parallel on different data streams
 - Multiprocessors including Intel Core2Dual, etc., fall into this category

Shared Memory Architecture: SMP/DSM (1/8):



- In such architecture, memory (centralized or distributed) can be directly accessed by different CPUs in the system
- Communication between programs/threads occurs implicitly via memory instructions (e.g., loads and stores)
- A natural extension of uni-processor model
 - Shared data are location transparent






Shared Memory Architecture: SMP/DSM (2/8):

SMP: Symmetric (Shared-Memory) Multiprocessors

- Intuitive parallel processing evolution in 1980s; several processors shared a single physical memory connected by a shared bus
 - **Symmetric shared memory** -- each processor has the same relationship to one single shared memory
 - Scalability limited by interconnection network






Shared Memory Architecture: SMP/DSM (3/8): SMP: Symmetric (Shared-Memory) Multiprocessors

- Characteristics

- UMA
 - Uniform memory access latency from any processor
- Symmetric (beyond UMA)
 - Equal access to computer's resources
 - Equal PE computing power
- Connected by bus or crossbar switch
- Shared-everything architecture
 - Centralized shared memory, shared I/O, same OSes
- Cache coherence protocol is needed if caches exist
 - Usually implemented by hardware




Shared Memory Architecture: SMP/DSM (4/8):

SMP: Symmetric (Shared-Memory) Multiprocessors

- Interconnection:
 - Bus
 - simple to design but cannot scale due to bus contention
 - Crossbar switch
 - Costly mechanism to avoid contention on different memory module (e.g., CPU1, CPU2, CPU3 all access different memory modules)
 - Whatever, necessary attribute is to offer equal PE-to-mem accesses
 - Difficult to scale to very large sizes:

2 to 32/128 processors today
High cache coherence overhead when scales




Examples of SMP Machines

- IBM eServer p5 590 16-way server
- IBM Netfinity 5500 M20, 4-way
- Sun Enterprise 10000, 64-way
- Sun Enterprise 6000 24-way
- Sun Fire V40z 8-way server
- Dell PowerEdge Model 6650, 4-way
- Dell PowerEdge 8450, 8-way,
- HP ProLiant DL580 G4 server, 4-way
- Fujitsu PRIMEPOWER 2500, 64-way
- (up to 128-way)
- Fujitsu Siemens Primergy K400, 4-way
- Unisys ES7000 Model 550, 32-way




Shared Memory Architecture: SMP/DSM (5/8): DSM: Distributed Shared-Memory Multiprocessors

- DSM: Distributed Shared-Memory Multiprocessors
 - Memory is distributed among the nodes, and all node-memory pairs are interconnected by a network
 - Memory access can be either local or remote




Shared Memory Architecture: SMP/DSM (6/8):

DSM: Distributed Shared-Memory Multiprocessors



- NUMA:
 - non-uniform memory accesses
 - All CPUs can read/write any part of the distributed memory, but
 - access time depends on location of a data word in memory
 - Local memory access
 - When local cache misses, and local memory hits
 - Shorter access latency than in UMA
 - Remote memory access
 - When local memory misses
 - Longer access latency than in UMA





Shared Memory Architecture: SMP/DSM (7/8): DSM: Distributed Shared-Memory Multiprocessors

- Benefits:
 - A cost-effective way to reduce memory traffic on the interconnection network:
 - By making most of the accesses to local memory
 - Improves application scalability
 - By exploiting data locality
- Cautions:
 - Memory access latency depends on **distance between processor and memory** – software must keep data close to the processor for scalable performance
 - Communications between processors becomes more complex – must be handled carefully by the architecture

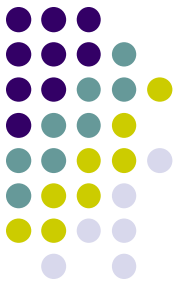
A decorative graphic in the bottom right corner consisting of a grid of colored dots in shades of purple, teal, yellow, and light blue, arranged in a pattern that tapers to the right.

The diagram illustrates two memory architectures on a horizontal axis representing machine scale, from 2 to "A few hundred processors".


Centralized Memory UMA: On the left, representing a smaller scale (2 to 32 processors), multiple processors (P_1, \dots, P_n) are connected to a central "Interconnection network" (yellow box). Each processor is connected to a local cache (labeled with a dollar sign '\$') which in turn connects to the network. The network is connected to a shared memory pool (dashed box labeled "Mem") which is connected to the network. This is labeled "Centralized Memory UMA" in red.

Distributed Memory (NUMA): On the right, representing a larger scale ("A few hundred processors"), each processor (P_1, \dots, P_n) is connected to a local cache (labeled with a dollar sign '\$') which is connected to a local memory module (labeled "Mem"). Each local memory module is connected to a local "Interconnection network" (yellow box). These local networks are connected to a larger, shared "Interconnection network" (yellow box). This is labeled "Distributed Memory (NUMA)" in red.

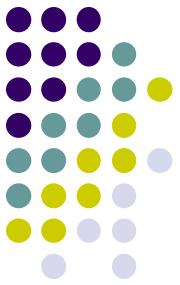
Shared Memory Communication (1/2):




- In shared memory systems, all processors share the same address space (main memory)
 - Address locations can be used to communicate implicitly via memory instructions (e.g. load and store)
- Small-scale systems usually centralize memory; result is a UMA

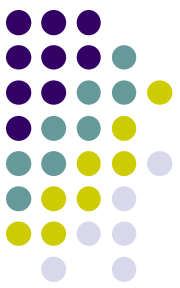


Shared Memory Communication (2/2):



- Large-scale multiprocessors must use **multiple memories that are physically distributed among the processors**; NUMA (small-scale multiprocessors? SMP is possible)
 - For communicating data among processors:
 - Physically separate memories can be addressed as one logically shared address space
 - The same physical address on two processors refers to the same location in memory (e.g., store A in P₁, load A in P_n)

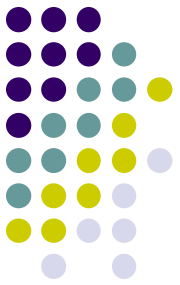




Cache Coherence in Shared Memory Architecture (1/19): Basic Schemes for Enforcing Coherence

- Programs on multiple processors may have copies of the same data in several caches
 - Most SMPs use **HW schemes** to maintain data coherence in the caches
 - Accessing data in memory/remote cache may take much longer time than accessing local cache
- Techniques to improve **data locality**
 - **Migration** - data can be moved to a local cache and used there in a transparent fashion
 - Reduces both
 - latency to access shared data that is allocated remotely and
 - bandwidth demand on the shared memory
 - **Replication** – for shared data being simultaneously read, since caches make a copy of data in local cache
 - Reduces both
 - latency of access and
 - contention for read shared data

Cache Coherence in Shared Memory Architecture (2/19):

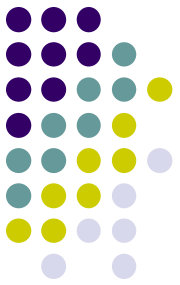


- Coherence and consistency
 - Coherence defines **values** returned by a read
 - Consistency determines **when** a written value will be returned by a read
- Cache Coherence problem
 - Processors see different values for X at time = 3

Write-through cache

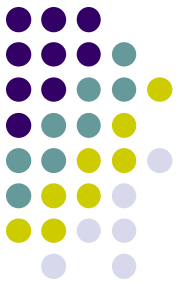
Time	Event	Cache contents for CPU A	Cache contents for CPU B	Memory contents for location X
0				1
1	CPU A reads X	1		1
2	CPU B reads X	1	1	1
3	CPU A stores 0 into X	0	1	0

Cache Coherence in Shared Memory Architecture (3/19):



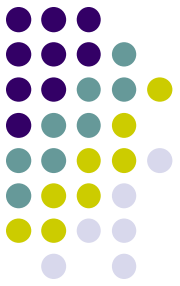
- Defining Coherent Memory System
 - 1. **Preserve Program Order:** A read by processor P to location X that follows a write by P to X, with no writes of X by another processor occurring between the write and the read by P, always returns the value written by P
 - 2. **Coherent view of memory:** Read by a processor to location X that follows a write by another processor to X returns the written value if the read and write are sufficiently separated in time and no other writes to X occur between the two accesses
 - 3. **Write serialization:** 2 writes to same location by any 2 processors are seen in the same order by all processors
 - If not, a processor could keep value 1 since saw as last write
 - For example, if the values 1 and then 2 are written to a location, processors can never read the value of the location as 2 and then later read it as 1

Cache Coherence in Shared Memory Architecture (4/19):




- Cache coherence protocols maintain coherence for multiple processors
- 2 Classes of Cache Coherence Protocols
 - Bus based -- **Snooping**
 - Every cache with a copy of data has a copy of sharing status of block, but no centralized state is kept
 - All caches are accessible via some broadcast medium (a bus or switch)
 - All cache controllers monitor or snoop on the bus to determine whether or not they have a copy of a block that is requested on the bus
 - **Directory based**
 - The sharing status of a block of physical memory is kept in just one location

Cache Coherence in Shared Memory Architecture (5/19):



- Example Write-Back Snoopy Protocol
 - Invalidation protocol, write-back cache
 - Snoops every address on bus
 - If it has a dirty copy of requested block, provides that block in response to the read request and aborts the memory access
 - Each memory block is in one state (but not to be indicated):
 - Clean in all caches and up-to-date in memory (Shared in cache)
 - OR Dirty in exactly one cache (Exclusive in cache)
 - OR not in any caches
 - Each cache block is in one state (track these):
 - Shared : block can be read
 - OR Exclusive : cache has only copy, it is writeable, and dirty
 - OR Invalid : block contains no data
 - Read misses: caches with that block respond
 - Writes to shared blocks are treated as misses (since cause same bus activities; simplifies design)

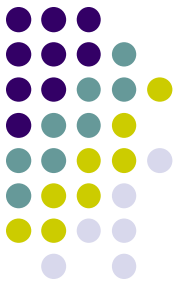
An Example of Invalidation Protocol



Processor Activity	Bus Activity	Contents of CPU A's Cache	Contents of CPU B's Cache	Contents of Memory Location X
				0

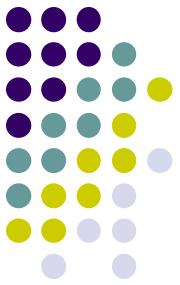
- An example of an invalidation protocol working on a snooping bus for a single cache block (X) with write-back caches.

An Example of Invalidation Protocol




Processor Activity	Bus Activity	Contents of CPU A's Cache	Contents of CPU B's Cache	Contents of Memory Location X
				0
CPU A reads X	Cache miss for X	0		0

An Example of Invalidation Protocol



Processor Activity	Bus Activity	Contents of CPU A's Cache	Contents of CPU B's Cache	Contents of Memory Location X
				0
CPU A reads X	Cache miss for X	0		0
CPU B reads X	Cache miss for X	0	0	0

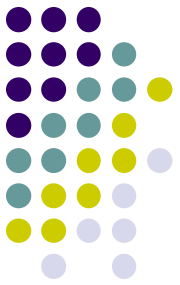
An Example of Invalidation Protocol



Processor Activity	Bus Activity	Contents of CPU A's Cache	Contents of CPU B's Cache	Contents of Memory Location X
				0
CPU A reads X	Cache miss for X	0		0
CPU B reads X	Cache miss for X	0	0	0
CPU A writes a 1 to X	Invalidation for X	1		0

- For a write, we require that the writing processor have exclusive access, preventing any other processor from being able to write simultaneously.

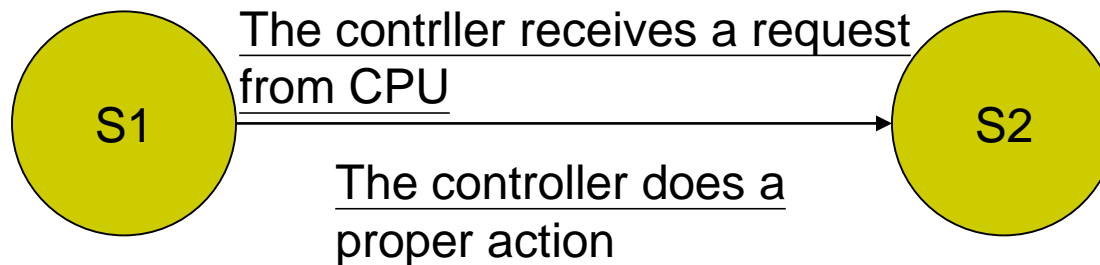
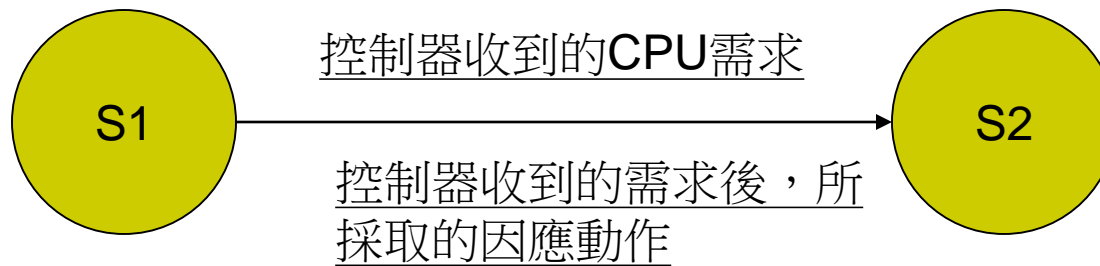
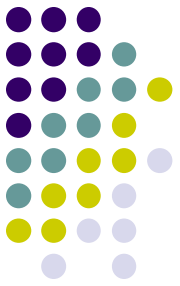
An Example of Invalidation Protocol




Processor Activity	Bus Activity	Contents of CPU A's Cache	Contents of CPU B's Cache	Contents of Memory Location X
				0
CPU A reads X	Cache miss for X	0		0
CPU B reads X	Cache miss for X	0	0	0
CPU A writes a 1 to X	Invalidation for X	1		0
CPU B reads X	Cache miss for X	1	1	1

- When the second miss by B occurs, CPU A responds with the value and cancels the response from memory. In addition, both the contents of B's cache and the memory contents of X are updated.

狀態機的狀態改變 (Finite State Machine: FSM)





Cache Coherence in Shared Memory Architecture (6/19):

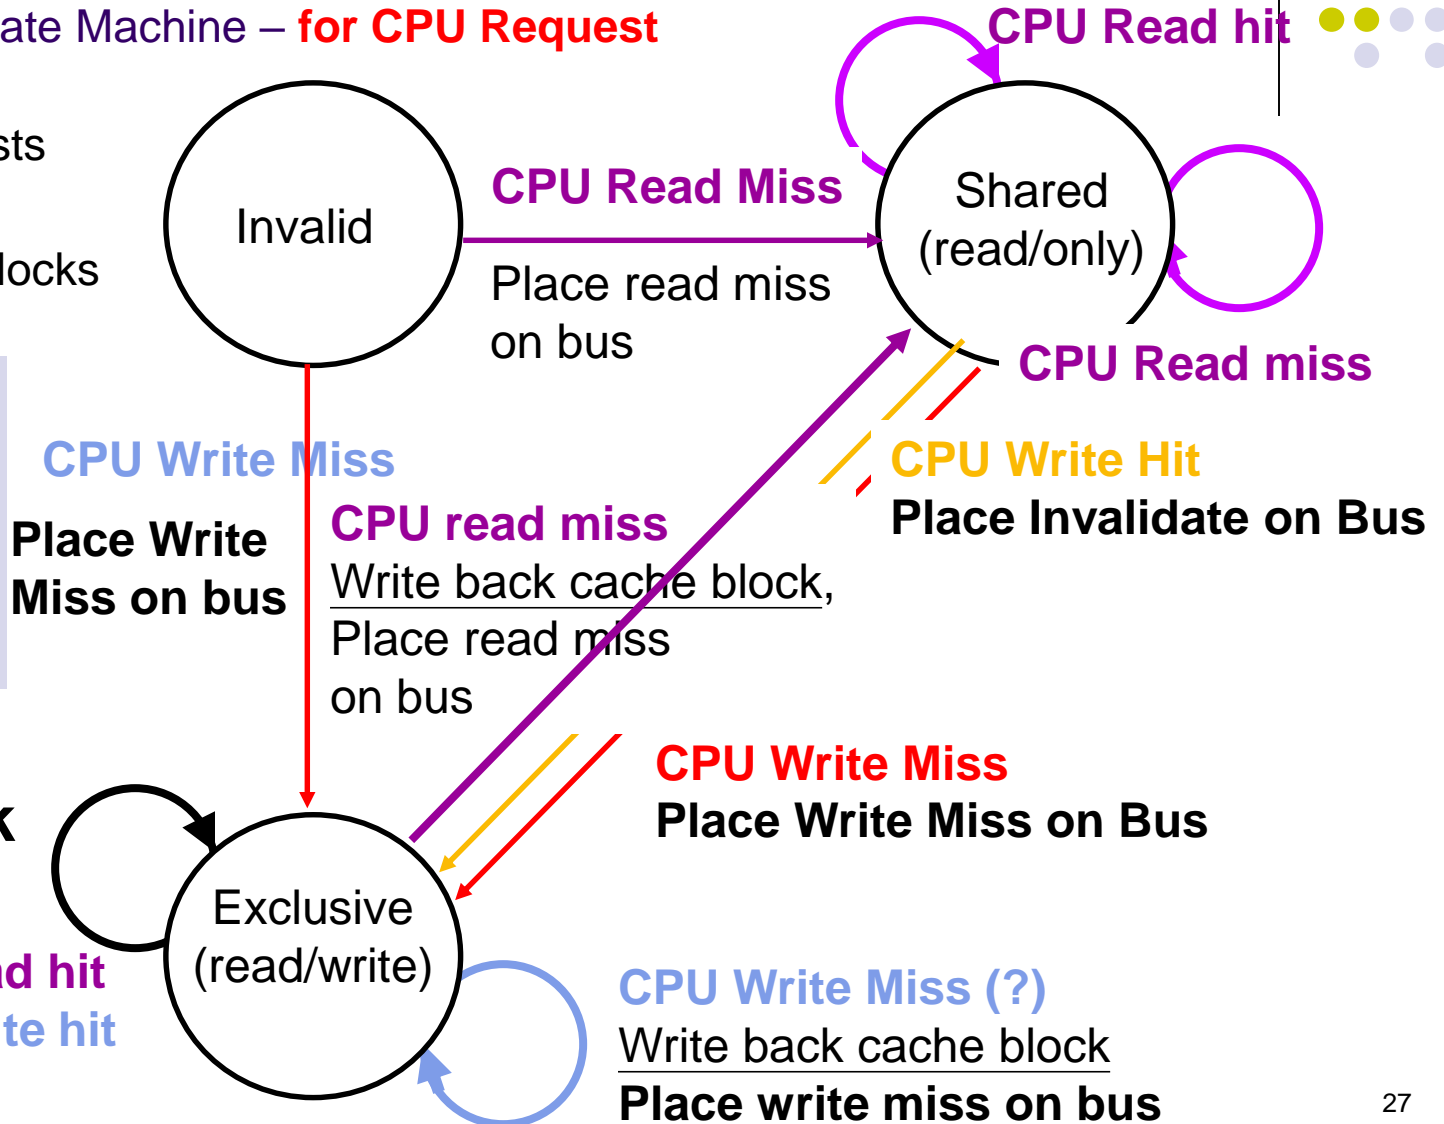
Write-Back State Machine – **for CPU Request**

- State machine for **CPU** requests for **each** **cache block**
- Non-resident blocks invalid

Cache controller幫每個Cache block維護一個對應的狀態機。
Cache controller maintains a finite state machine for each cache block.

Cache Block State

CPU read hit
CPU write hit

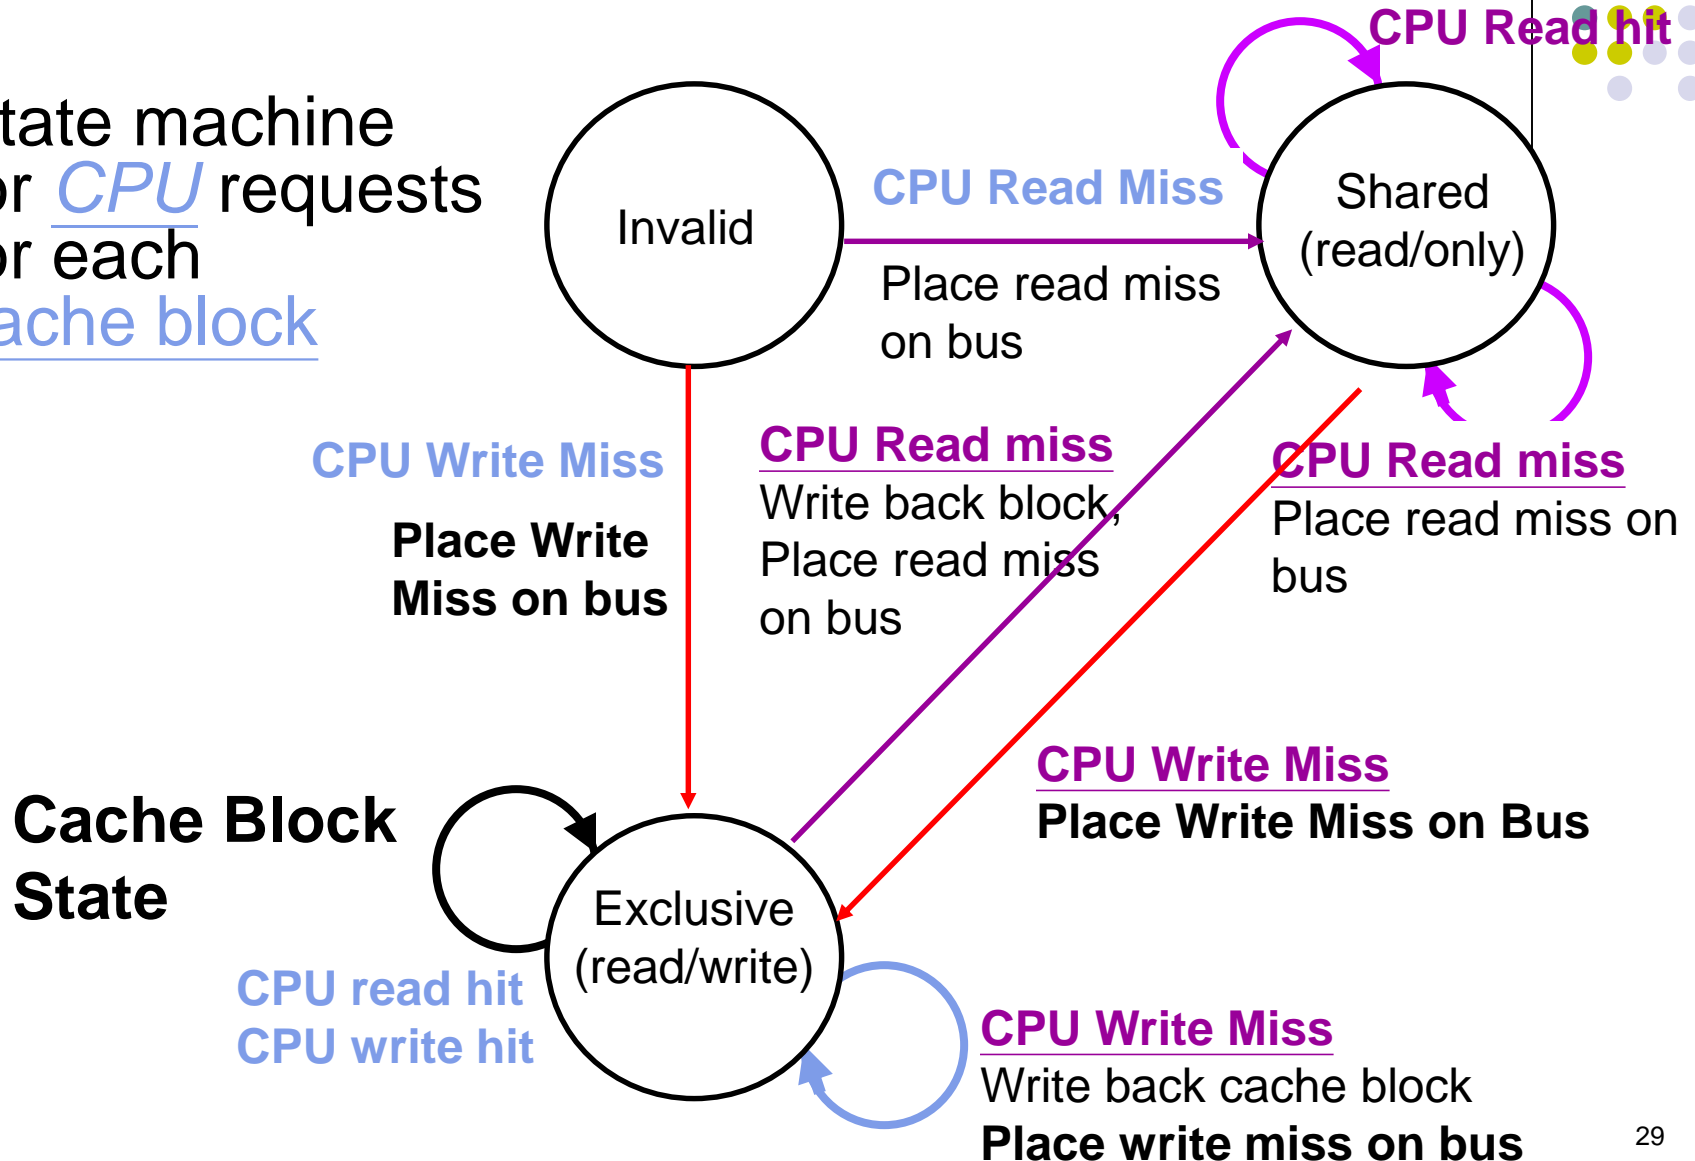


-
- The diagram illustrates the states and transitions of a memory consistency protocol. It features three circular nodes: Invalid, Shared (read/only), and Exclusive (read/write). Transitions are represented by arrows with associated labels:
- Invalid to Shared (read/only):** Write miss for this block (cyan arrow)
 - Shared (read/only) to Invalid:** Invalidate for this block (grey arrow)
 - Exclusive (read/write) to Invalid:** Write miss for this block (cyan arrow)
 - Exclusive (read/write) to Shared (read/only):** Read miss for this block, Write Back Block; (abort memory access) (teal arrow)
- ```
graph TD; Invalid((Invalid)); Shared((Shared (read/only))); Exclusive((Exclusive (read/write))); Invalid -- "Write miss for this block" --> Shared; Shared -- "Invalidate for this block" --> Invalid; Exclusive -- "Write miss for this block" --> Invalid; Exclusive -- "Read miss for this block, Write Back Block; (abort memory access)" --> Shared;
```

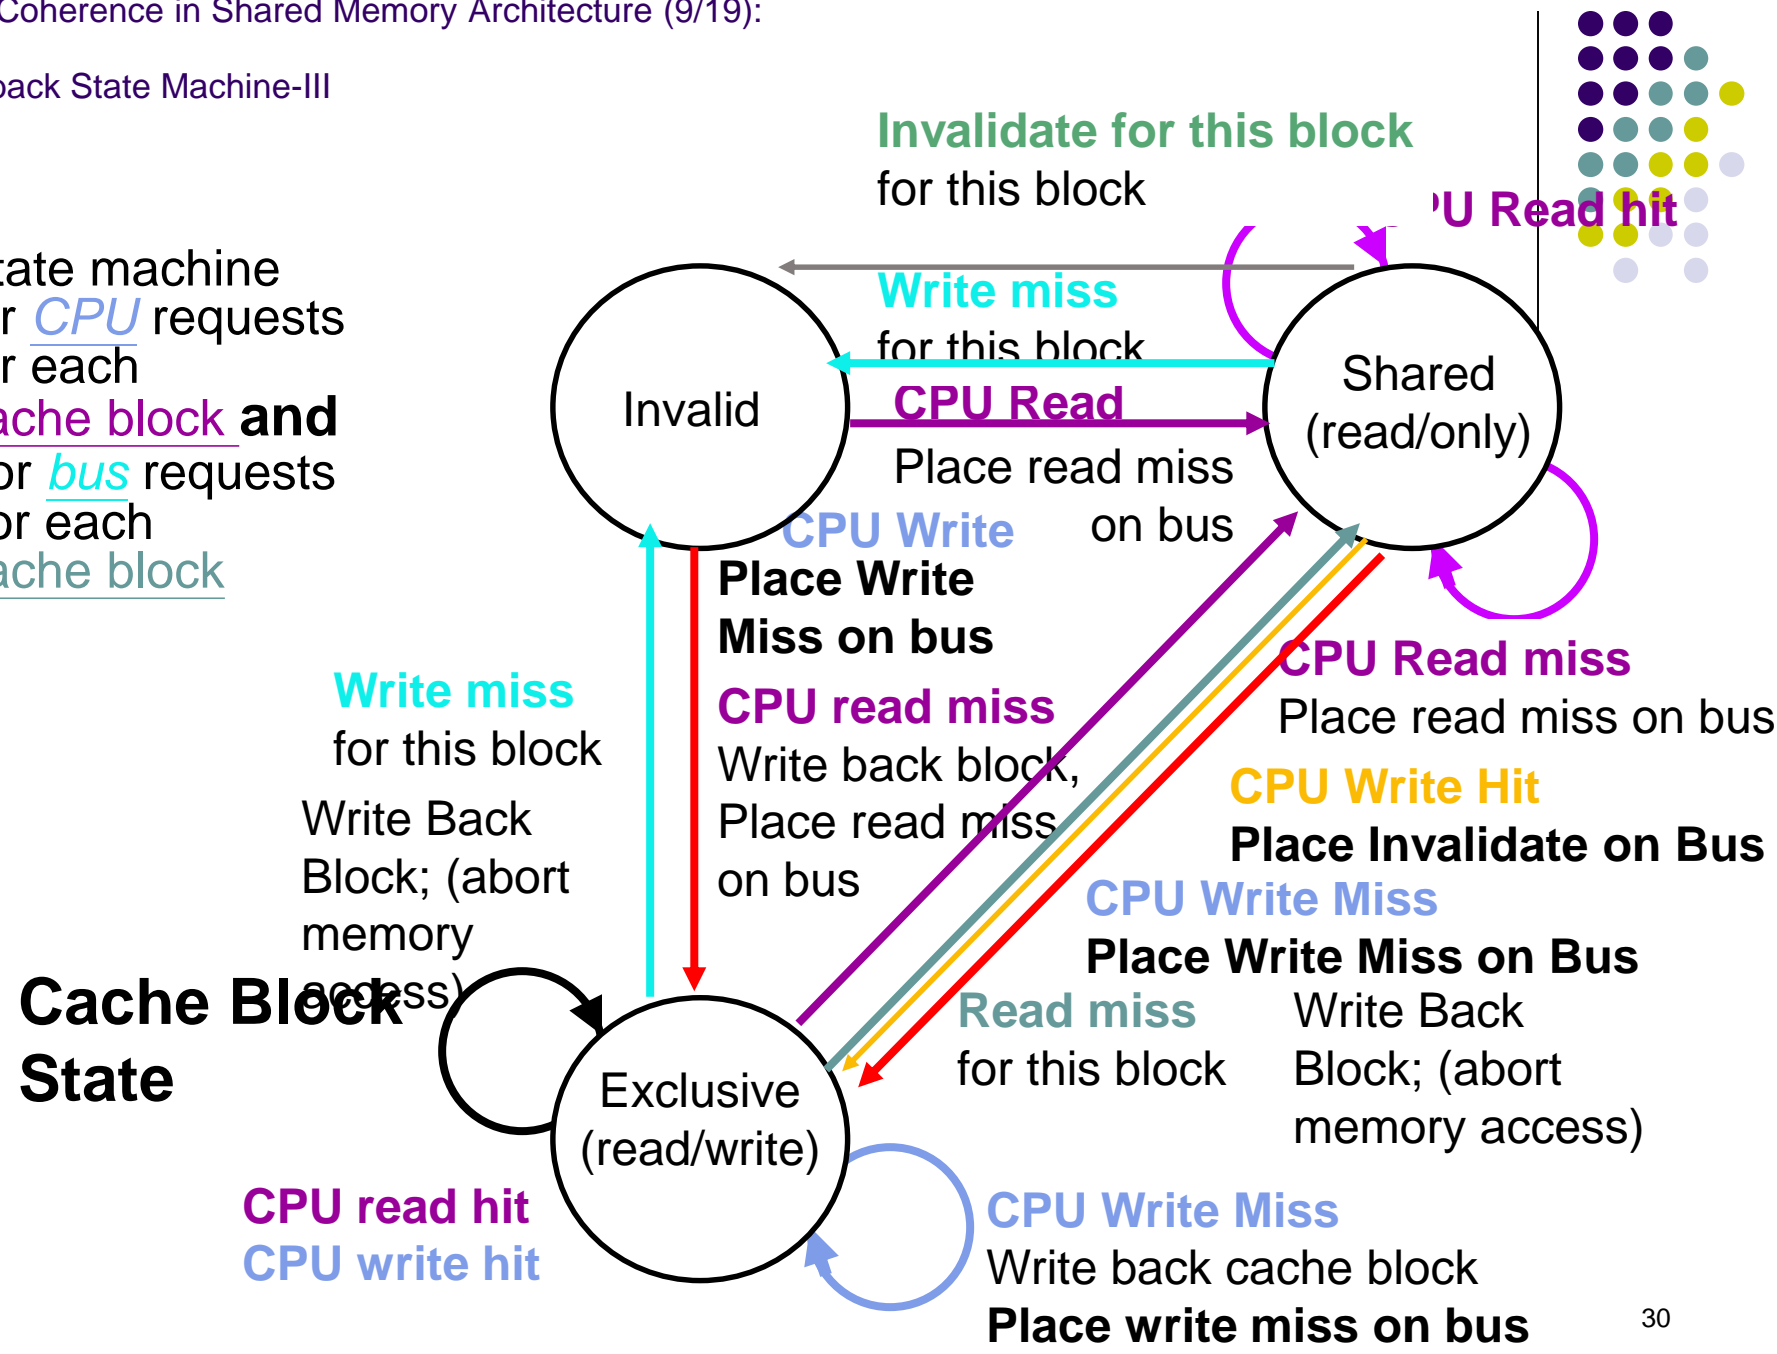
# Cache Coherence in Shared Memory Architecture (8/19): Block-replacement



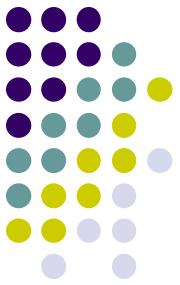
- State machine for CPU requests for each cache block



- State machine for CPU requests for each cache block and for bus requests for each cache block



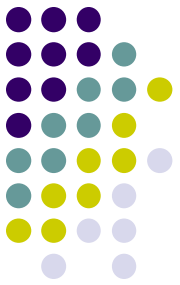
# Example



| step               | P1    |      |       | P2    |      |       | Bus    |       |      |       |  | Memory |       |
|--------------------|-------|------|-------|-------|------|-------|--------|-------|------|-------|--|--------|-------|
|                    | State | Addr | Value | State | Addr | Value | Action | Proc. | Addr | Value |  | Addr   | Value |
| P1 Write 10 to A1  |       |      |       |       |      |       |        |       |      |       |  |        |       |
| P1: Read A1        |       |      |       |       |      |       |        |       |      |       |  |        |       |
| P2: Read A1        |       |      |       |       |      |       |        |       |      |       |  |        |       |
|                    |       |      |       |       |      |       |        |       |      |       |  |        |       |
|                    |       |      |       |       |      |       |        |       |      |       |  |        |       |
| P2: Write 20 to A1 |       |      |       |       |      |       |        |       |      |       |  |        |       |
| P2: Write 40 to A2 |       |      |       |       |      |       |        |       |      |       |  |        |       |
|                    |       |      |       |       |      |       |        |       |      |       |  |        |       |

Assumes A1 and A2 map to same cache block,  
initial cache state is invalid

# Example

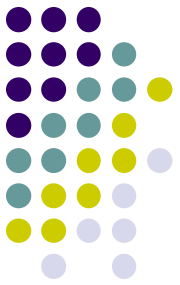


|                    |  | P1           |           |           | P2    |      |       | Bus         |      |      | Memory |     |      |
|--------------------|--|--------------|-----------|-----------|-------|------|-------|-------------|------|------|--------|-----|------|
| step               |  | State        | Addr      | Value     | State | Addr | Value | Action      | Proc | Addr | Value  | Add | Valu |
| P1 Write 10 to A1  |  | <u>Excl.</u> | <u>A1</u> | <u>10</u> |       |      |       | <u>WrMs</u> | P1   | A1   |        |     |      |
| P1: Read A1        |  |              |           |           |       |      |       |             |      |      |        |     |      |
| P2: Read A1        |  |              |           |           |       |      |       |             |      |      |        |     |      |
|                    |  |              |           |           |       |      |       |             |      |      |        |     |      |
|                    |  |              |           |           |       |      |       |             |      |      |        |     |      |
| P2: Write 20 to A1 |  |              |           |           |       |      |       |             |      |      |        |     |      |
| P2: Write 40 to A2 |  |              |           |           |       |      |       |             |      |      |        |     |      |
|                    |  |              |           |           |       |      |       |             |      |      |        |     |      |

Assumes A1 and A2 map to same cache block



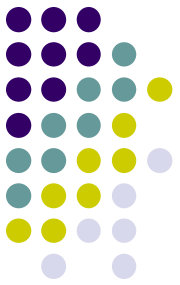
# Example



| step               | P1           |           |           | P2    |      |       | Bus         |       |      |       | Memory |       |
|--------------------|--------------|-----------|-----------|-------|------|-------|-------------|-------|------|-------|--------|-------|
|                    | State        | Addr      | Value     | State | Addr | Value | Action      | Proc. | Addr | Value | Addr   | Value |
| P1 Write 10 to A1  | <u>Excl.</u> | <u>A1</u> | <u>10</u> |       |      |       | <u>WrMs</u> | P1    | A1   |       |        |       |
| P1: Read A1        | Excl.        | A1        | 10        |       |      |       |             |       |      |       |        |       |
| P2: Read A1        |              |           |           |       |      |       |             |       |      |       |        |       |
|                    |              |           |           |       |      |       |             |       |      |       |        |       |
|                    |              |           |           |       |      |       |             |       |      |       |        |       |
| P2: Write 20 to A1 |              |           |           |       |      |       |             |       |      |       |        |       |
| P2: Write 40 to A2 |              |           |           |       |      |       |             |       |      |       |        |       |
|                    |              |           |           |       |      |       |             |       |      |       |        |       |

Assumes A1 and A2 map to same cache block

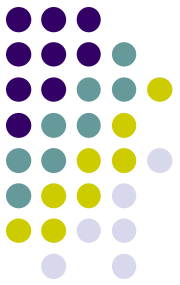
# Example



|                           | <i>P1</i>    |             |              | <i>P2</i>    |             |              | <i>Bus</i>    |             |             | <i>Memory</i> |             |              |
|---------------------------|--------------|-------------|--------------|--------------|-------------|--------------|---------------|-------------|-------------|---------------|-------------|--------------|
| <i>step</i>               | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>Action</i> | <i>Proc</i> | <i>Addr</i> | <i>Value</i>  | <i>Addr</i> | <i>Value</i> |
| <b>P1 Write 10 to A1</b>  | <u>Excl.</u> | <u>A1</u>   | <u>10</u>    |              |             |              | <u>WrMs</u>   | P1          | A1          |               |             |              |
| <b>P1: Read A1</b>        | Excl.        | A1          | 10           |              |             |              |               |             |             |               |             |              |
| <b>P2: Read A1</b>        |              |             |              | <u>Shar.</u> | <u>A1</u>   |              | <u>RdMs</u>   | P2          | A1          |               |             |              |
|                           | <u>Shar.</u> | A1          | 10           | Shar.        | A1          | <u>10</u>    | <u>WrBk</u>   | P1          | A1          | 10            | A1          | <u>10</u>    |
| <b>P2: Write 20 to A1</b> |              |             |              |              |             |              |               |             |             |               |             |              |
| <b>P2: Write 40 to A2</b> |              |             |              |              |             |              |               |             |             |               |             |              |
|                           |              |             |              |              |             |              |               |             |             |               |             |              |

Assumes A1 and A2 map to same cache block

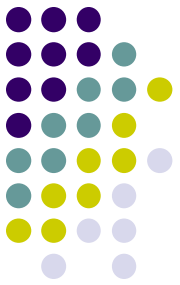
# Example



| step               | P1           |           |           | P2           |           |           | Bus         |      |      | Memory |           |           |
|--------------------|--------------|-----------|-----------|--------------|-----------|-----------|-------------|------|------|--------|-----------|-----------|
|                    | State        | Addr      | Value     | State        | Addr      | Value     | Action      | Proc | Addr | Value  | Add       | Value     |
| P1 Write 10 to A1  | <u>Excl.</u> | <u>A1</u> | <u>10</u> |              |           |           | <u>WrMs</u> | P1   | A1   |        |           |           |
| P1: Read A1        | Excl.        | A1        | 10        |              |           |           |             |      |      |        |           |           |
| P2: Read A1        |              |           |           | <u>Shar.</u> | <u>A1</u> |           | <u>RdMs</u> | P2   | A1   |        |           |           |
|                    | <u>Shar.</u> | A1        | 10        | Shar.        | A1        | <u>10</u> | <u>WrBk</u> | P1   | A1   | 10     | <u>A1</u> | <u>10</u> |
| P2: Write 20 to A1 | <u>Inv.</u>  |           |           | <u>Excl.</u> | A1        | <u>20</u> | <u>Inv</u>  | P2   | A1   |        | A1        | 10        |
| P2: Write 40 to A2 |              |           |           |              |           |           |             |      |      |        |           |           |
|                    |              |           |           |              |           |           |             |      |      |        |           |           |
|                    |              |           |           |              |           |           |             |      |      |        |           |           |

Assumes A1 and A2 map to same cache block


# Example



|                           | <i>P1</i>    |             |              | <i>P2</i>    |             |              | <i>Bus</i>    |             |             |              | <i>Memory</i> |              |
|---------------------------|--------------|-------------|--------------|--------------|-------------|--------------|---------------|-------------|-------------|--------------|---------------|--------------|
| <i>step</i>               | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>Action</i> | <i>Proc</i> | <i>Addr</i> | <i>Value</i> | <i>Add</i>    | <i>Value</i> |
| <b>P1 Write 10 to A1</b>  | <u>Excl.</u> | <u>A1</u>   | <u>10</u>    |              |             |              | <u>WrMs</u>   | P1          | A1          |              |               |              |
| <b>P1: Read A1</b>        | Excl.        | A1          | 10           |              |             |              |               |             |             |              |               |              |
| <b>P2: Read A1</b>        |              |             |              | <u>Shar.</u> | <u>A1</u>   |              | <u>RdMs</u>   | P2          | A1          |              |               |              |
|                           | <u>Shar.</u> | A1          | 10           | Shar.        | A1          | <u>10</u>    | <u>WrBk</u>   | P1          | A1          | 10           | <u>A1</u>     | <u>10</u>    |
| <b>P2: Write 20 to A1</b> | <u>Inv.</u>  |             |              | <u>Excl.</u> | A1          | <u>20</u>    | <u>Inv</u>    | P2          | A1          |              | A1            | 10           |
| <b>P2: Write 40 to A2</b> |              |             |              |              |             |              | <u>WrBk</u>   | P2          | A1          | 20           | <u>A1</u>     | <u>20</u>    |
|                           |              |             |              | Excl.        | <u>A2</u>   | <u>40</u>    | <u>WrMs</u>   | P2          | A2          |              | A1            | 20           |
|                           |              |             |              |              |             |              |               |             |             |              |               |              |

Assumes A1 and A2 map to same cache block,  
but A1 != A2

# Cache Coherence in Shared Memory Architecture (10/19): Coherency Misses



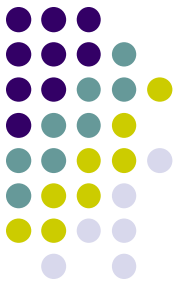
為了確保coherence之下，導致原已讀取的block被丟棄，重新再讀取。

## ● 4th C: *coherence miss*

### ● Joins Compulsory, Capacity, Conflict

- 1. **True sharing misses** arise from the communication of data through the cache coherence mechanism
  - The first write by a processor to a shared cache block causes an invalidation to establish ownership of that block.
  - When another processor attempts to read a **modified word** in that cache block, a miss occurs and the resultant block is transferred.
  - Miss would still occur if block size were 1 word
- 2. **False sharing misses** when a block is invalidated because **some words in the block, other than the one being read, is written into**
  - Invalidation does not cause a new value to be communicated, but only causes an extra cache miss
  - Block is shared, but no word in block is actually shared  
⇒ miss would not occur if block size were 1 word

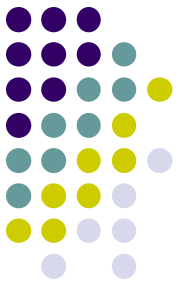
# Cache Coherence in Shared Memory Architecture (11/19): Coherency Misses



- Example: True vs. False Sharing vs. Hit?
  - Assume x1 and x2 in same cache block
    - P1 and P2 both read x1 and x2 before

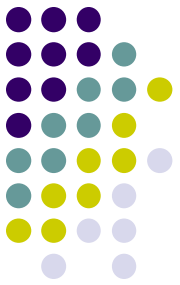
| Time | P1       | P2       | True, False, Hit? Why?                                     |
|------|----------|----------|------------------------------------------------------------|
| 1    | Write x1 |          | Write hit; place invalidate on bus;<br>invalidate x1 in P2 |
| 2    |          | Read x2  | <b>False miss; x1 irrelevant to P2</b>                     |
| 3    | Write x1 |          | Write hit; place invalidate on bus;<br>invalidate x1 in P2 |
| 4    |          | Write x2 | <b>False miss; invalidate x2 in P1</b>                     |
| 5    | Read x2  |          | <b>True miss</b>                                           |

# Cache Coherence in Shared Memory Architecture (12/19): Directory Protocol



- For larger-scale architectures, memory bandwidth also need to be scaled
  - Bus-based coherence **lacks scalability**
  - Scalable Approach: Directories
    - A directory keeps the state of every memory block that may be cached.
      - Information in the directory includes which caches have copies of the memory block, whether it is dirty, and so on.



# Cache Coherence in Shared Memory Architecture (13/19): Directory Protocol



- **Similar to Snoopy Protocol: Three states for each memory block**
  - **Shared**: One or more processors have the block cached, and the value in memory is up to date.
  - **Uncached**: No processor has a copy of the memory block.
  - **Exclusive**: Exactly one processor (owner) has a copy of the memory block, and it has written the block, so the memory copy is out of date.
- In addition to cache state, we must track **which processors have data** when in the shared state (usually **bit vector**, 1 if processor has copy)



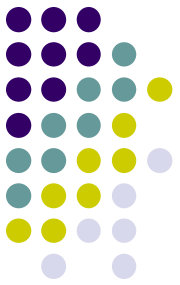
# Cache Coherence in Shared Memory Architecture (14/19): Directory Protocol




- k processors.
- With each block in memory:  
k presence-bits, 1 dirty-bit
- With each block in cache:  
1 valid bit, and 1 dirty (owner) bit

- Read from main memory by processor i:
  - If dirty-bit OFF then { read from main memory; turn p[i] ON; }
  - if dirty-bit ON then { recall line from dirty proc (cache state to shared); update memory; turn dirty-bit OFF; turn p[i] ON; supply recalled data to i; }
- Write to main memory by processor i:
  - If dirty-bit OFF then { supply data to i; send invalidations to all caches that have the block; turn dirty-bit ON; turn p[i] ON; ... }
- ...

# Cache Coherence in Shared Memory Architecture (15/19): Directory Protocol



- **Terms: Typically 3 processors involved**
  - **Local node** where a read/write **request** originates
  - **Home node** where the **memory location** and the **directory** entry of an address reside
  - **Remote node** has a **copy** of a cache block, whether exclusive or shared
- **Physical address space**
  - The high-order bits may provide the **node number**
  - The low-order bits provide the **offset** within the memory on that node



# Cache Coherence in Shared Memory Architecture (16/19): Directory Protocol

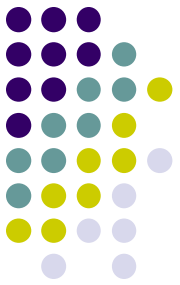
| Message Type       | Source         | Destination    | Message contents | Function of this message                                                                                             |
|--------------------|----------------|----------------|------------------|----------------------------------------------------------------------------------------------------------------------|
| Read miss          | local cache    | home directory | P, A             | Processor P has a read miss at address A; request data and make P a read sharer.                                     |
| Write miss         | local cache    | home directory | P, A             | Processor P has a write miss at address A; request data and make P the exclusive owner.                              |
| Invalidate         | local cache    | home directory | A                | Request to send invalidates to all remote caches that are caching the block at address A.                            |
| Invalidate         | home directory | remote cache   | A                | Invalidate a shared copy of data at address A.                                                                       |
| Fetch              | home directory | remote cache   | A                | Fetch the block at address A and send it to its home directory; change the state of A in the remote cache to shared. |
| Fetch / invalidate | home directory | remote cache   | A                | Fetch the block at address A and send it to its home directory; invalidate the block in the cache.                   |
| Data value reply   | home directory | local cache    | D                | Return a data value form the home memory.                                                                            |
| Data write back    | remote cache   | home directory | A, D             | Write back a data value for address A.                                                                               |

P = requesting processor number

A = requested address

D = data contents

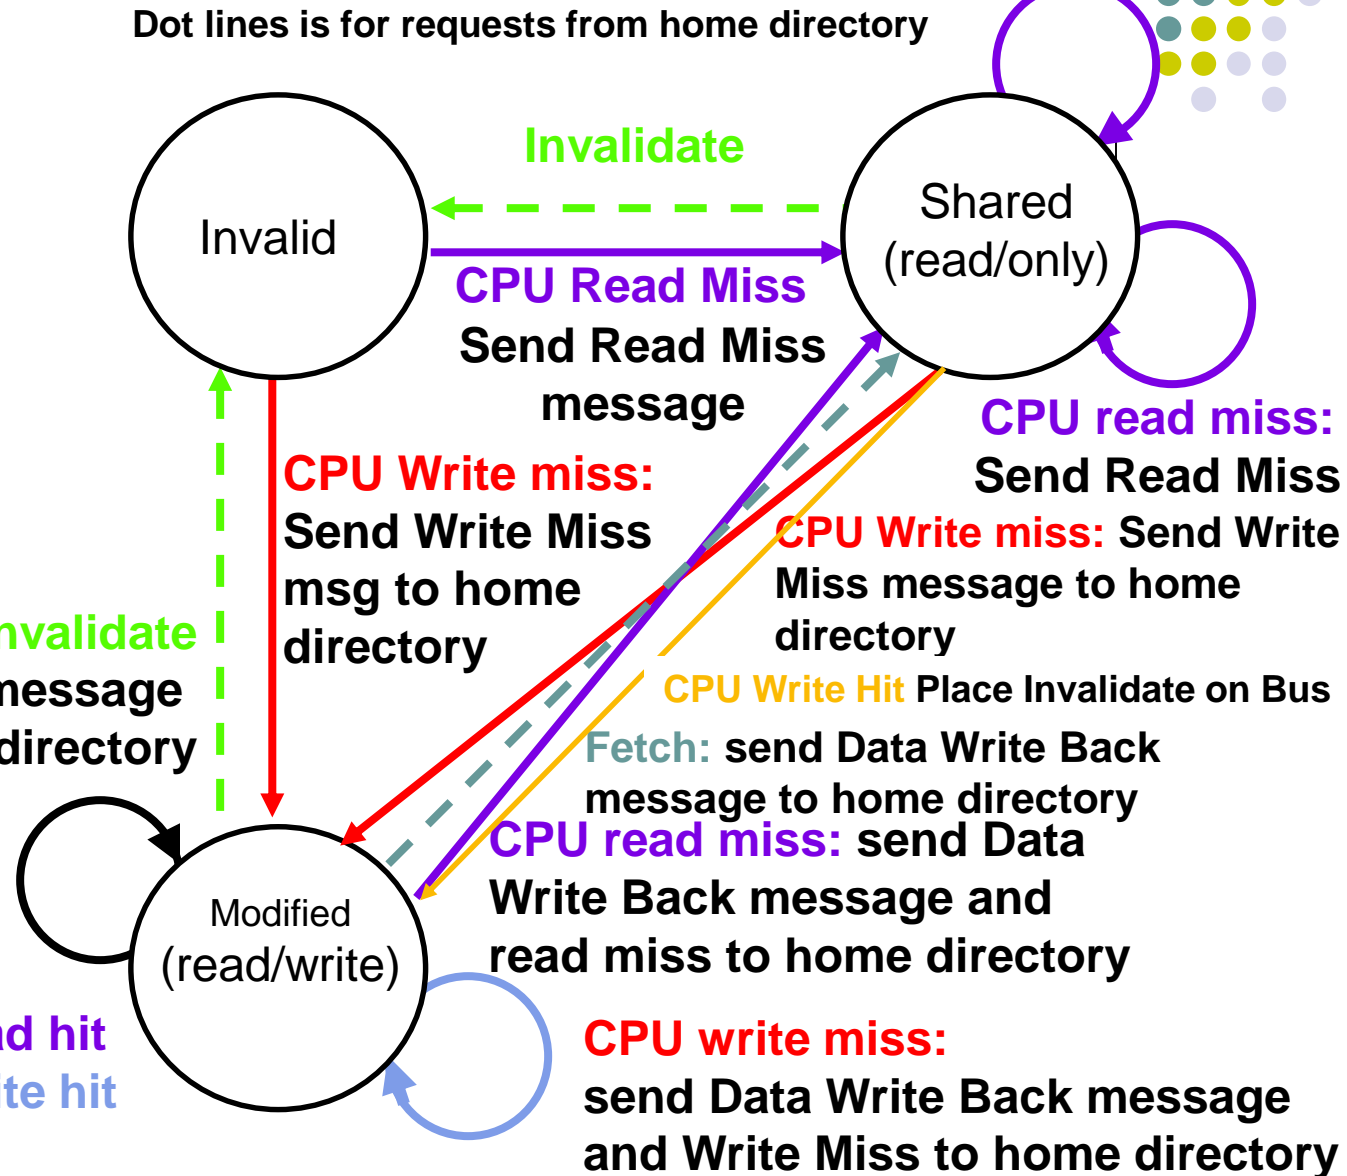
# Cache Coherence in Shared Memory Architecture (17/19): Directory Protocol



- Cache states identical to snoopy case; transactions very similar
- State transitions for a cache caused by **read misses**, **write misses**, **invalidates**, **data fetch** requests
- An individual cache generates **read miss**, **write miss**, and **invalidate** message to home directory
- Write misses that were broadcast on the bus for snooping ⇒ **explicit invalidate & data fetch requests**

# Lead hit

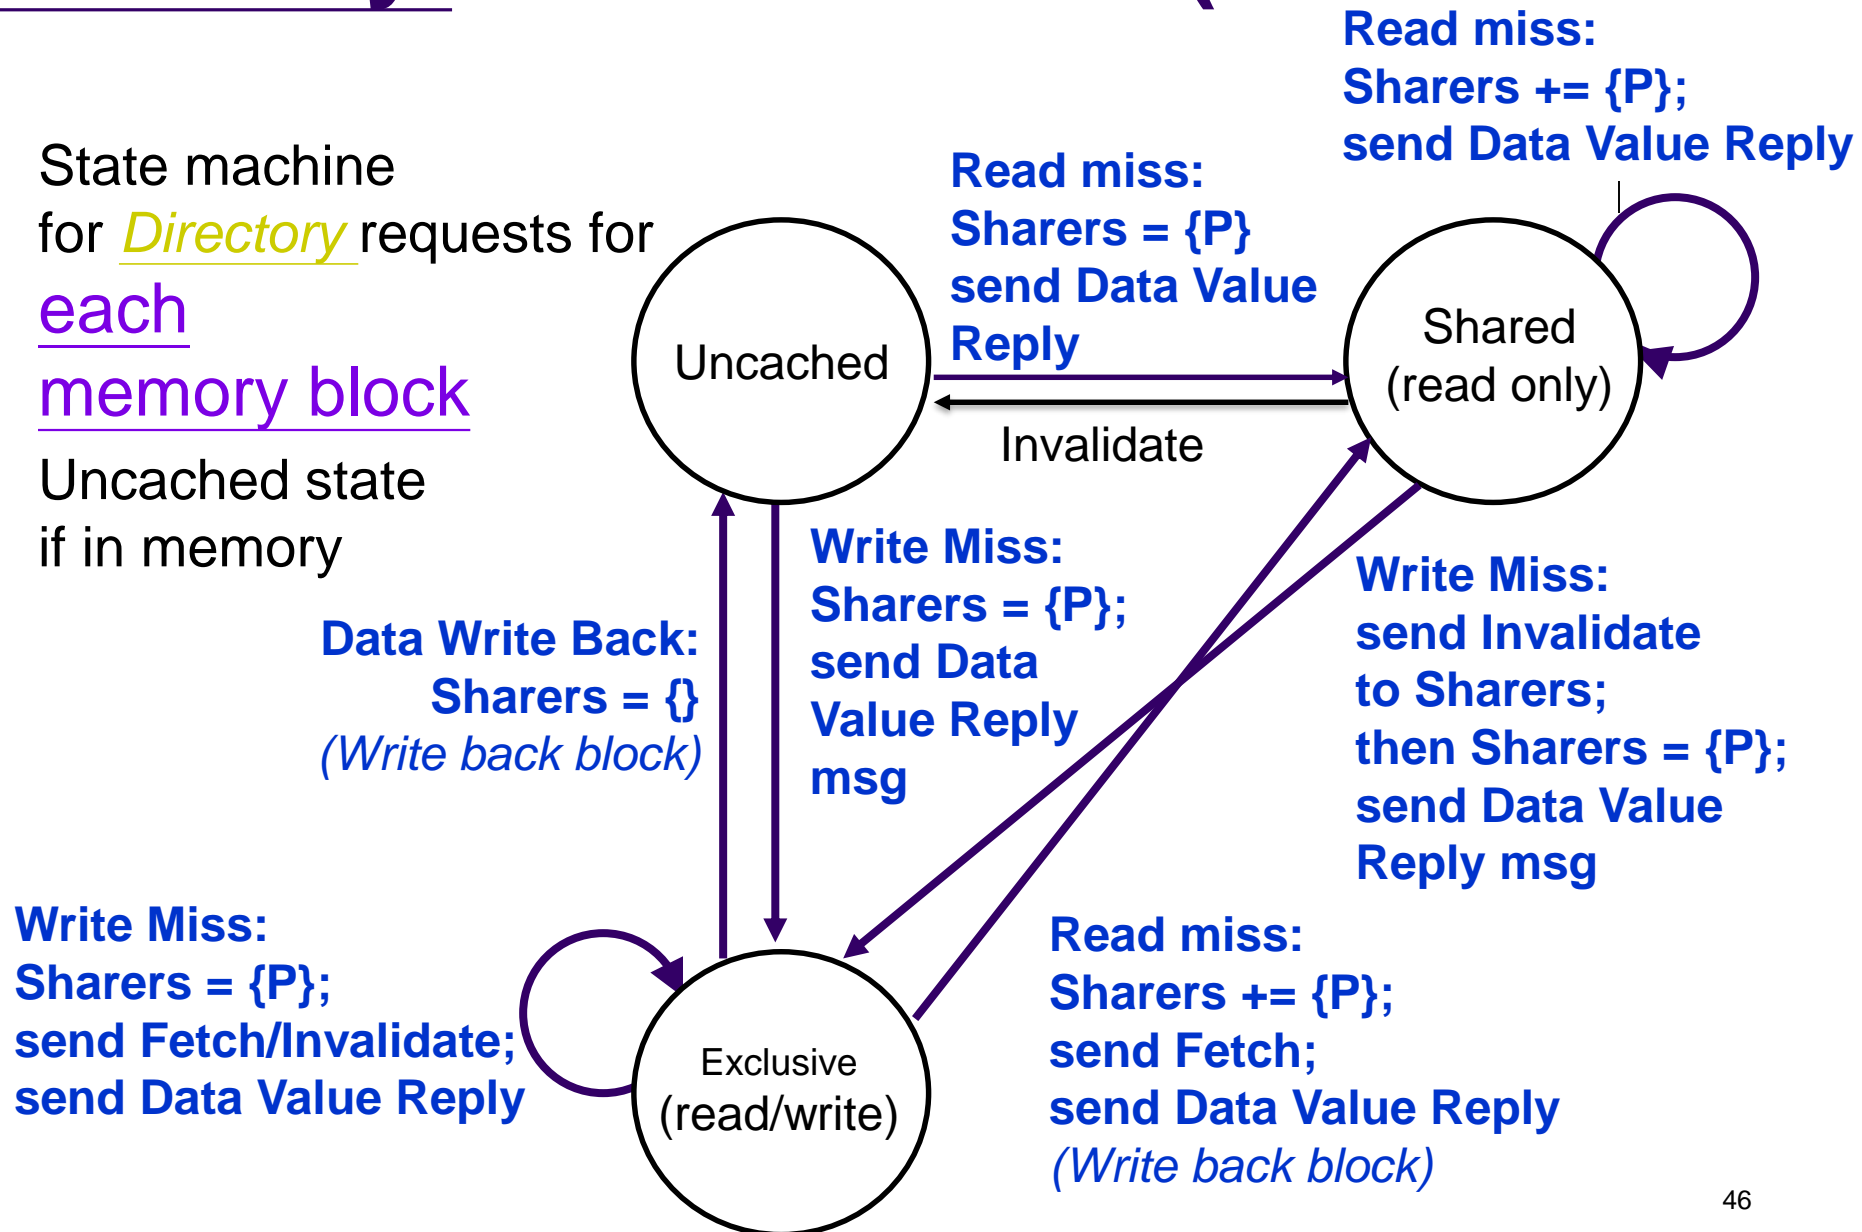
- Fetch/Invalidate**  
**send Data Write Back message**  
**to home directory**



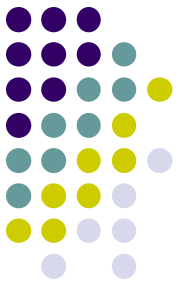
# Directory State Machine (19/19)



- State machine for Directory requests for each memory block
- Uncached state if in memory

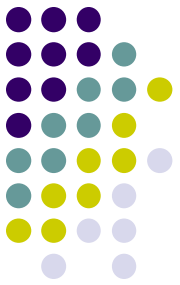


# State Transition Diagram for Directory



- Same states & structure as the transition diagram for an individual cache
- A message sent to a directory causes two different types of actions:
  - update the directory state
  - send messages to satisfy the request
- Unlike snoopy, the directory state indicates the state of all the cached copies of a memory block, rather than for a single cache block.
- The directory must track the set of processors that have a copy of a block; we use a set called **Sharers** to perform this function.

# Example




Processor 1    Processor 2    Interconnect    Directory    Memory

|                    | <i>P1</i>    |             |              | <i>P2</i>    |             |              | <i>Bus</i>    |              |             | <i>Directory</i> |             |              | <i>Memor</i>   |              |
|--------------------|--------------|-------------|--------------|--------------|-------------|--------------|---------------|--------------|-------------|------------------|-------------|--------------|----------------|--------------|
| <i>step</i>        | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>Action</i> | <i>Proc.</i> | <i>Addr</i> | <i>Value</i>     | <i>Addr</i> | <i>State</i> | <i>{Procs}</i> | <i>Value</i> |
| P1: Write 10 to A1 |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
| P1: Read A1        |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
| P2: Read A1        |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
| P2: Write 20 to A1 |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
| P2: Write 40 to A2 |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |                  |             |              |                |              |

A1 and A2 map to the same cache block



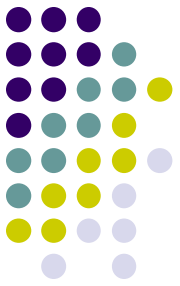


# Example

|                    | Processor 1  |             |              | Processor 2  |             |              | Interconnect  |              |             |              | Directory        |              |                | Memory       |
|--------------------|--------------|-------------|--------------|--------------|-------------|--------------|---------------|--------------|-------------|--------------|------------------|--------------|----------------|--------------|
|                    | <i>P1</i>    |             |              | <i>P2</i>    |             |              | <i>Bus</i>    |              |             |              | <i>Directory</i> |              |                | <i>Memor</i> |
| <i>step</i>        | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>Action</i> | <i>Proc.</i> | <i>Addr</i> | <i>Value</i> | <i>Addr</i>      | <i>State</i> | <i>{Procs}</i> | <i>Value</i> |
| P1: Write 10 to A1 |              |             |              |              |             |              | <i>WrMs</i>   | P1           | A1          |              | <i>A1</i>        | <i>Ex</i>    | <i>{P1}</i>    |              |
|                    | <i>Excl.</i> | <i>A1</i>   | <i>10</i>    |              |             |              | <i>DaRp</i>   | P1           | A1          | 0            |                  |              |                |              |
| P1: Read A1        |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
| P2: Read A1        |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
| P2: Write 20 to A1 |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
| P2: Write 40 to A2 |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |
|                    |              |             |              |              |             |              |               |              |             |              |                  |              |                |              |

A1 and A2 map to the same cache block


# Example



Processor 1    Processor 2    Interconnect    Directory    Memory

|                    |  | <i>P1</i>    |             |              | <i>P2</i>    |             |              | <i>Bus</i>    |              |             |              | <i>Directory</i> |              |                | <i>Memory</i> |
|--------------------|--|--------------|-------------|--------------|--------------|-------------|--------------|---------------|--------------|-------------|--------------|------------------|--------------|----------------|---------------|
| <i>step</i>        |  | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>State</i> | <i>Addr</i> | <i>Value</i> | <i>Action</i> | <i>Proc.</i> | <i>Addr</i> | <i>Value</i> | <i>Addr</i>      | <i>State</i> | <i>{Procs}</i> | <i>Value</i>  |
| P1 Write 10 to A1  |  |              |             |              |              |             |              | <u>WrMs</u>   | P1           | A1          |              | <u>A1</u>        | <u>Ex</u>    | <u>{P1}</u>    |               |
|                    |  | <u>Excl.</u> | <u>A1</u>   | <u>10</u>    |              |             |              | <u>DaRp</u>   | P1           | A1          | 0            |                  |              |                |               |
| P1: Read A1        |  | Excl.        | A1          | 10           |              |             |              |               |              |             |              | <u>A1</u>        | <u>Ex</u>    | <u>{P1}</u>    |               |
| P2: Read A1        |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |
| P2: Write 20 to A1 |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |
|                    |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |
| P2: Write 20 to    |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |
|                    |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |
| P2: Write 40 to    |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |
|                    |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |
|                    |  |              |             |              |              |             |              |               |              |             |              |                  |              |                |               |

# Example



# Example

Step3

Step5

Step1

Step2

Step4

Processor 1

Processor 2

Interconnect


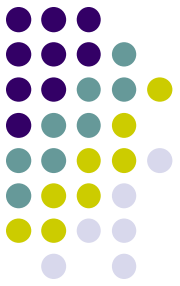
Directory

Memory

| step               | P1<br>State | P1<br>Addr | P1<br>Value | P2<br>State | P2<br>Addr | P2<br>Value | Bus<br>Action | Proc     | Addr     | Value | Directory<br>Addr | Directory<br>State | Directory<br>{Procs} | Memory<br>Value |
|--------------------|-------------|------------|-------------|-------------|------------|-------------|---------------|----------|----------|-------|-------------------|--------------------|----------------------|-----------------|
| P1 Write 10 to A1  | Excl.       | A1         | 10          |             |            |             | WrMs<br>DaRp  | P1<br>P1 | A1<br>A1 | 0     | A1                | Ex                 | {P1}                 |                 |
| P1: Read A1        | Excl.       | A1         | 10          |             |            |             |               |          |          |       | A1                | Ex                 | {P1}                 |                 |
| P2: Read A1        |             |            |             |             |            |             | RdMs<br>Ftch  | P2<br>P1 | A1<br>A1 |       |                   |                    |                      |                 |
|                    | Shar.       | A1         | 10          |             |            |             | DaRp          | P2       | A1       | 10    | A1                | Shar.              | {P1,P2}              | A1=10           |
| P2: Write 20 to A1 |             |            |             | Shar.       | A1         | 10          |               |          |          |       |                   |                    |                      |                 |
| P2: Write 40 to A2 |             |            |             |             |            |             |               |          |          |       |                   |                    |                      |                 |
|                    |             |            |             |             |            |             |               |          |          |       |                   |                    |                      |                 |
|                    |             |            |             |             |            |             |               |          |          |       |                   |                    |                      |                 |

A1 and A2 map to the same cache block

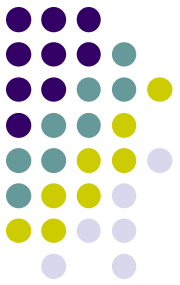
# Example



| step               |  | P1           |           |           | P2           |      |           | Bus           |      |      | Directory |           |              | Memory         |              |
|--------------------|--|--------------|-----------|-----------|--------------|------|-----------|---------------|------|------|-----------|-----------|--------------|----------------|--------------|
|                    |  | State        | Addr      | Value     | State        | Addr | Value     | Action        | Proc | Addr | Value     | Addr      | State        | {Procs}        | Value        |
| P1 Write 10 to A1  |  |              |           |           |              |      |           | <u>WrMs</u>   | P1   | A1   |           | <u>A1</u> | <u>Ex</u>    | <u>{P1}</u>    |              |
|                    |  | <u>Excl.</u> | <u>A1</u> | <u>10</u> |              |      |           | <u>DaRp</u>   | P1   | A1   | 0         |           |              |                |              |
| P1: Read A1        |  | Excl.        | A1        | 10        |              |      |           |               |      |      |           |           |              |                |              |
| P2: Read A1        |  |              |           |           |              |      |           | <u>RdMs</u>   | P2   | A1   |           |           |              |                |              |
|                    |  | <u>Shar.</u> | A1        | 10        |              |      |           | <u>Ftch</u>   | P1   | A1   | 10        |           |              |                | <u>A1=10</u> |
|                    |  |              |           |           | Shar.        | A1   | <u>10</u> | <u>DaRp</u>   | P2   | A1   | 10        | A1        | <u>Shar.</u> | <u>{P1,P2}</u> |              |
| P2: Write 20 to A1 |  |              |           |           |              |      |           | <u>Inval.</u> | P1   | A1   |           |           |              |                |              |
|                    |  | <u>Inv.</u>  |           |           | <u>Excl.</u> | A1   | <u>20</u> |               |      |      |           | A1        | <u>Excl.</u> | <u>{P2}</u>    |              |
| P2: Write 40 to A2 |  |              |           |           |              |      |           |               |      |      |           |           |              |                |              |
|                    |  |              |           |           |              |      |           |               |      |      |           |           |              |                |              |
|                    |  |              |           |           |              |      |           |               |      |      |           |           |              |                |              |

A1 and A2 map to the same cache block

# Example



# Example

Step1

Step3

Processor 1

Processor 2

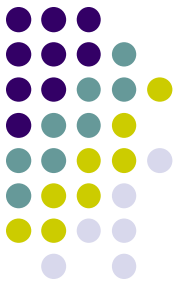
Interconnect

Directory

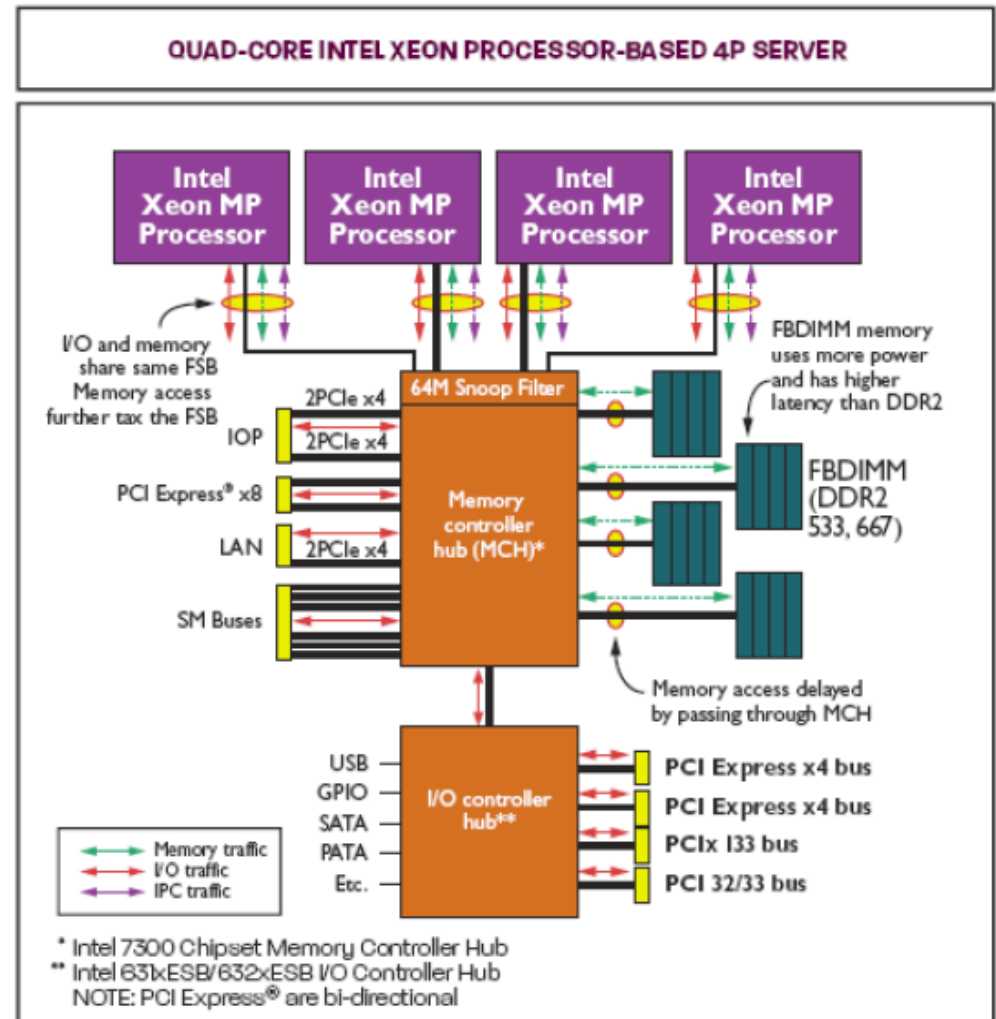
Memory

|                    |       |      |       |       |      |       |        |       |      |       |           |       |         |       |
|--------------------|-------|------|-------|-------|------|-------|--------|-------|------|-------|-----------|-------|---------|-------|
|                    | P1    |      |       | P2    |      |       | Bus    |       |      |       | Directory |       | Memory  |       |
| step               | State | Addr | Value | State | Addr | Value | Action | Proc. | Addr | Value | Addr      | State | {Procs} | Value |
| P1 Write 10 to A1  |       |      |       |       |      |       | WrMs   | P1    | A1   |       | A1        | Ex    | {P1}    |       |
|                    | Excl. | A1   | 10    |       |      |       | DaRp   | P1    | A1   | 0     |           |       |         |       |
| P1: Read A1        | Excl. | A1   | 10    |       |      |       |        |       |      |       |           |       |         |       |
| P2: Read A1        |       |      |       |       |      |       | RdMs   | P2    | A1   |       |           |       |         |       |
|                    | Shar. | A1   | 10    |       |      |       | Ftch   | P1    | A1   | 10    |           |       |         | A1=10 |
|                    |       |      |       | Shar. | A1   | 10    | DaRp   | P2    | A1   | 10    | A1        | Shar. | {P1,P2} |       |
| P2: Write 20 to A1 |       |      |       | Excl. | A1   | 20    |        |       |      |       |           |       |         |       |
|                    | Inv.  |      |       |       |      |       | Inval. | P1    | A1   |       | A1        | Excl. | {P2}    |       |
| P2: Write 40 to A2 |       |      |       |       |      |       | WrMs   | P2    | A2   |       | A2        | Excl. | {P2}    |       |
|                    |       |      |       |       |      |       | WrBk   | P2    | A1   | 20    | A1        | Unca. | { }     | A1=20 |
|                    |       |      |       | Excl. | A2   | 40    | DaRp   | P2    | A2   | 0     |           |       |         |       |

# Cache Coherence in Shared Memory Architecture



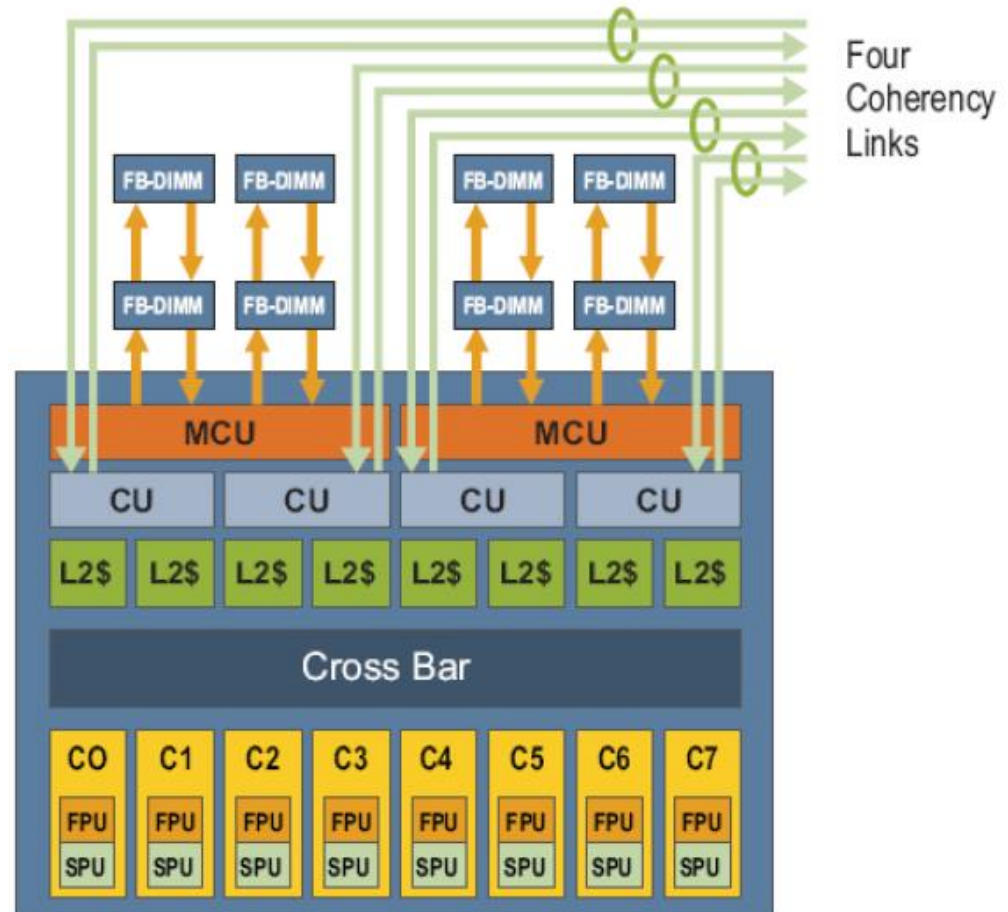
- Intel QuadCore SMP
  - Snoopy protocol
  - Snoop Filter cache



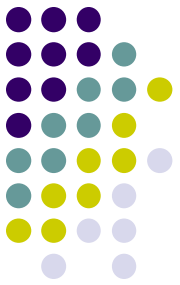
# Cache Coherence in Shared Memory Architecture



- Example:
  - UltraSPARC T2 Plus
    - NUMA architecture
    - 4 Coherency Units (CUs).
    - cache coherence (snoopy)



# Distributed Memory Architecture (1/2):




- Each CPU has its own memory view and associated memory
  - CPUs exchange data between their respective memory spaces when required, via **message passing**
  - Hence called **NORMA**: NO Remote-Memory Access
- User must be aware of sender/receiver of communicated data, and will have to move or distribute these data explicitly when needed
- Distributed-memory MIMD systems exhibit a large variety in their connecting topology (due to better flexibility)




# Distributed Memory Architecture (2/2):




- NORMA machine examples:
  - MPP (massively parallel processors), cluster, ...
- Popular topologies:
  - (Fat) tree, mesh/torus, cube, ...



2D Mesh

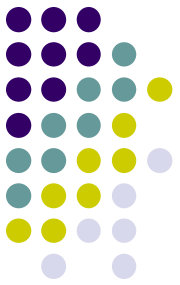


3D Torus



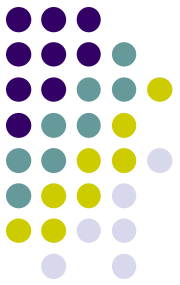
Hypercube (d=4)

# Distributed Memory Communication: Message-Passing (1/2):

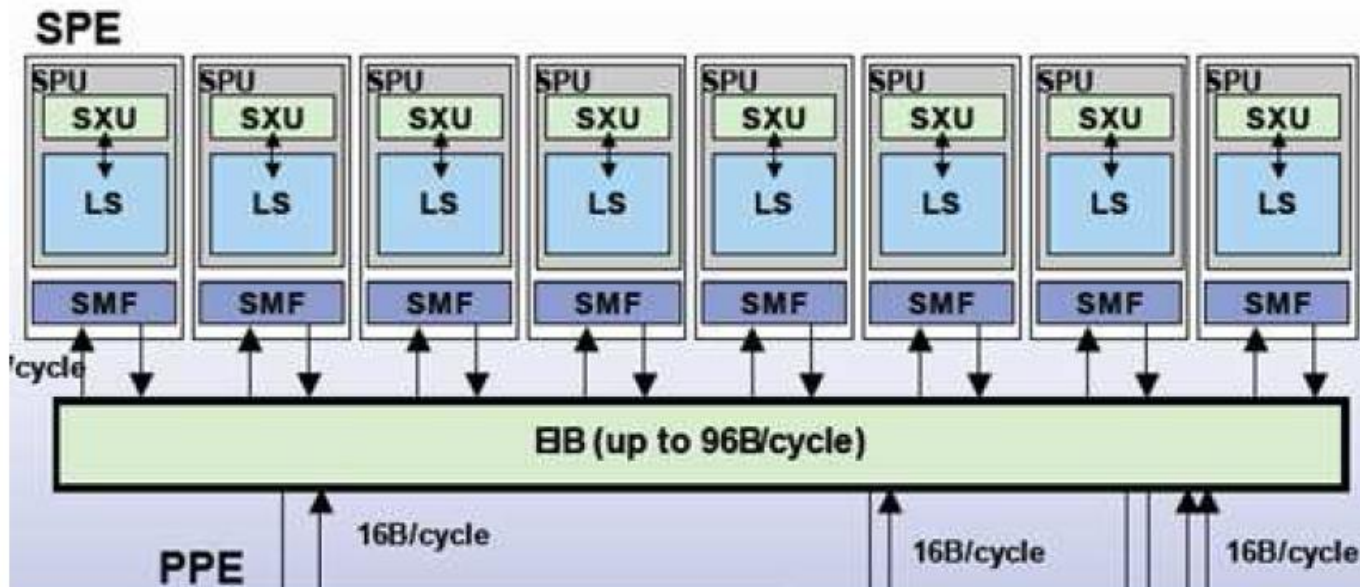


- Message passing multiprocessors
  - Multiprocessors are with multiple address spaces
    - communication of data is done by explicitly passing messages among the processors
- Message passing communication
  - Communications occurs by sending/receiving messages
    - request and deliver data just as with the network protocols
  - Can be *synchronous* or *asynchronous*
    - **Synchronous**: the initiating processor sends a request and waits until the reply is returned before continuing
    - **Asynchronous**: the initiating processor sends a request, yet does not wait for the reply; it immediately continues execution

# Distributed Memory Communication: Message-Passing (2/2):



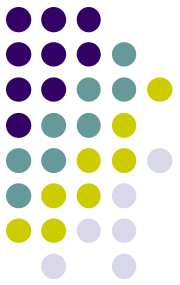
- Example:
  - IBM Cell SPEs
    - each SPE is an independent processor which can run its own programs
    - A shared, coherent memory and a rich set of DMA commands provide for communications between all Cell processing elements





## Shared Memory vs. Distributed Memory (1/3): Advantages of Shared Memory

- Compatibility
  - Same communication mechanism works for all machines. Can have standardized programming interface (as OpenMP has proposed)
- Ease of programming
  - Obvious when communication patterns are complex or vary dynamically during execution. Similar advantage simplifies compiler design
  - Programmer allowed to focus attention only on performance-critical accesses
- Lower overhead for communication and better use of bandwidth when communicating small items
  - Arises from the implicit nature of communication and use of memory mapping to implement protection in hardware, rather than through I/O system
- Ability to use hardware-controlled caching to reduce remote communications
  - By supporting automatic caching of all data, both shared and private.



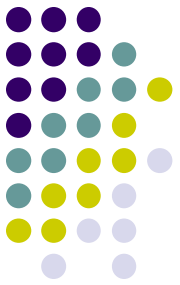
## Shared Memory vs. Distributed Memory (2/3): Advantages of Distributed Memory

- Hardware can be simpler
  - in comparison with a shared memory implementation that supports coherent caching of remote data
- Communication is explicitly written in programs via sending messages
- Explicit communication captures programmer attention on this costly aspect of parallel computation
  - sometimes leading to improved structure in a multiprocessor program.
- Synchronization is naturally associated with sending/receiving messages
  - reducing possibility for errors introduced by incorrect synchronization
- Easier to use sender-initiated communication
  - may have some performance advantage



## Shared Memory vs. Distributed Memory (3/3):

- Shared memory architecture is more powerful than distributed memory architecture
  - Cache coherence hardware moves data automatically
  - Can run both shared-memory and message-passing programs
- Distributed memory architecture is simpler and may take considerably less execution time than shared memory architecture
  - Software is responsible for moving data
  - Runs only message-passing programs
- In general,
  - Parallelizing programs for shared-memory is easier and takes less time, but it takes further efforts to both debug and scale the performance
  - Parallelizing programs for distributed-memory is harder and takes longer time, but programs are better behaved, and higher performance can be gained



# Summary

- Shared memory
- Cache coherence problem
- Snooping protocol
- Directory protocol
- Distributed memory